



Convnets in TensorFlow

CS 20: TensorFlow for Deep Learning Research

Lecture 7

2/7/2017

Agenda

Convolutions without training

Convnet with MNIST!!!

`tf.layers`



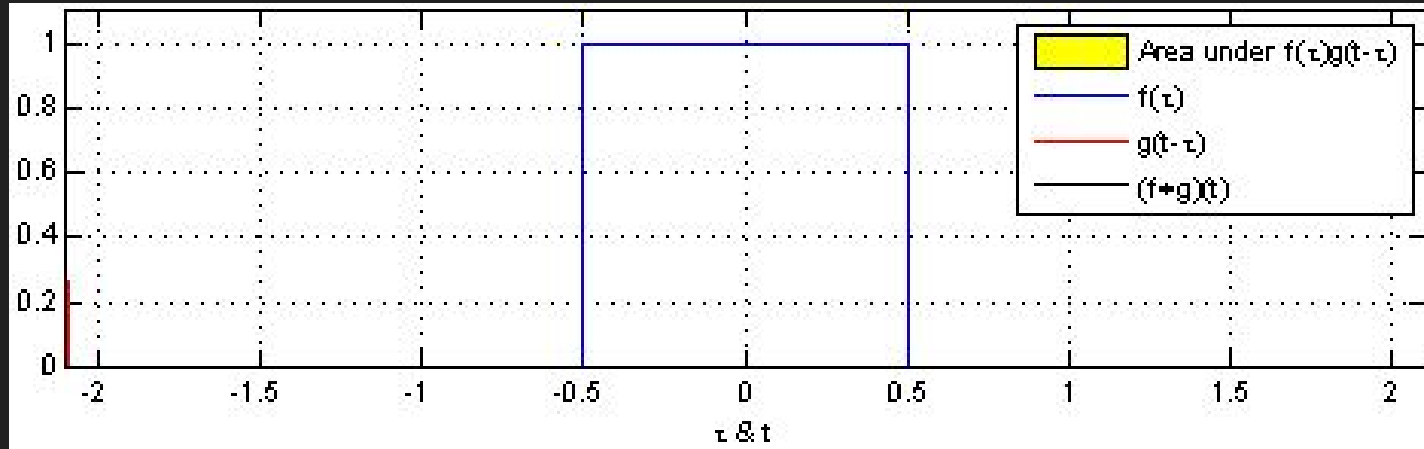


Understanding convolutions

Convolutions in math and physics

a function derived from two given functions by integration that expresses how the shape of one is modified by the other

Convolutions in math and physics



Convolutions in math and physics

How an input is transformed by a kernel*

*also called filter/feature map

Convolutions in machine learning

We can use one single convolutional layer to modify a certain image

Convolutions in machine learning

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

Kernel for blurring

0.0625	0.125	0.0625
0.125	0.25	0.125
0.0625	0.125	0.0625

Matrix multiplication of this kernel with
a 3 x 3 patch of an image is a weighted sum
of neighboring pixels
=> blurring effect

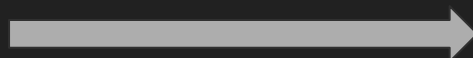
Convolution without training



input

Kernel for blurring

0.0625	0.125	0.0625
0.125	0.25	0.125
0.0625	0.125	0.0625



`tf.nn.conv2d`



output

Convolutions in TensorFlow

We can use one single convolutional layer to modify a certain image

```
tf.nn.conv2d(  
    input,  
    filter,  
    strides,  
    padding,  
    use_cudnn_on_gpu=True,  
    data_format='NHWC',  
    dilations=[1, 1, 1, 1],  
    name=None  
)
```

Convolutions in TensorFlow

We can use one single convolutional layer to modify a certain image

```
tf.nn.conv2d(  
    input,           Batch size (N) x Height (H) x Width (W) x Channels (C)  
    filter,          Height x Width x Input Channels x Output Channels  
    strides,         4 element 1-D tensor, strides in each direction  
    padding,         'SAME' or 'VALID'  
    use_cudnn_on_gpu=True,  
    data_format='NHWC',  
    dilations=[1, 1, 1, 1],  
    name=None  
)
```

Convolutions in TensorFlow

We can use one single convolutional layer to modify a certain image

```
tf.nn.conv2d(  
    image,  
    kernel,  
    strides=[1, 3, 3, 1],  
    padding='SAME',  
)
```



Some basic kernels



input



sharpen



edge



top sobel



emboss

See `kernels.py` and `07_run_kernels.py`

Some basic kernels



input



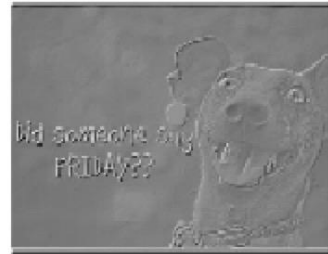
sharpen



edge



top sobel



emboss



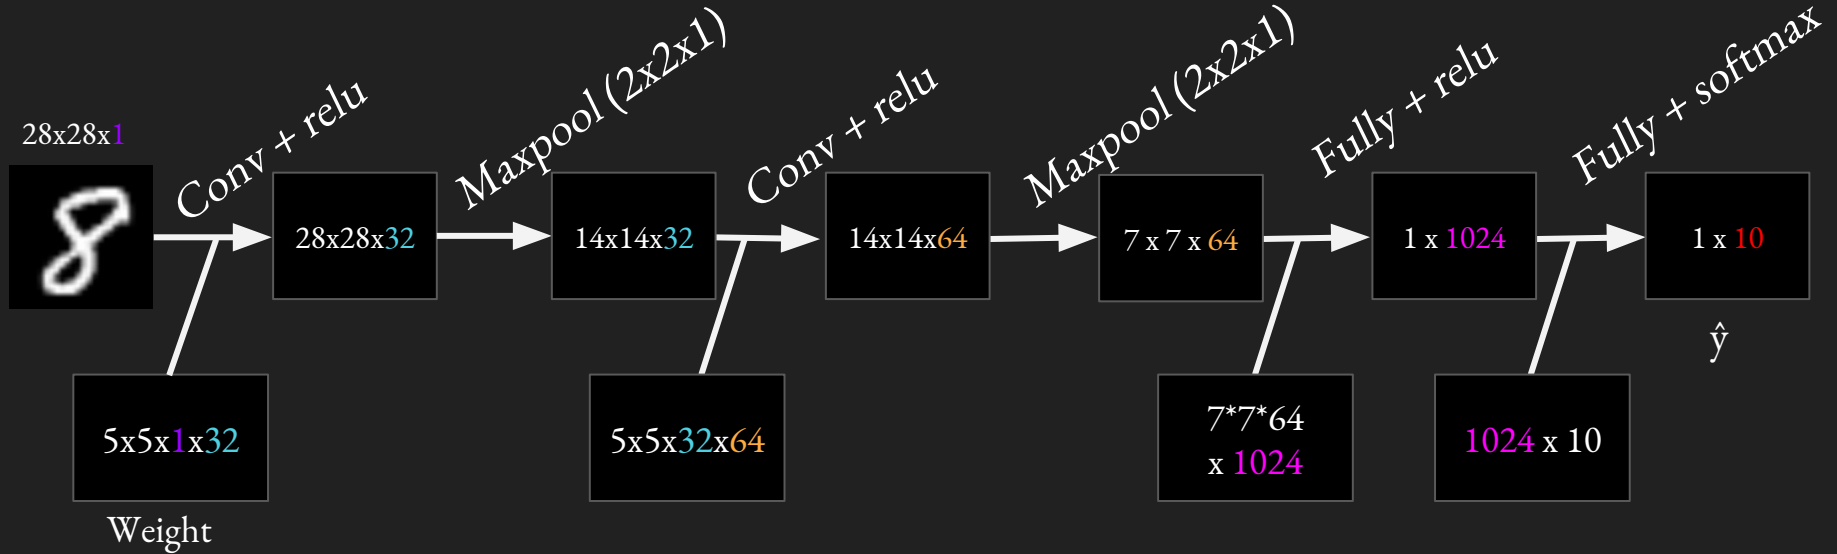
Convolutions in machine learning

Don't hard-code the values of your kernels.
Learn the optimal kernels through training!



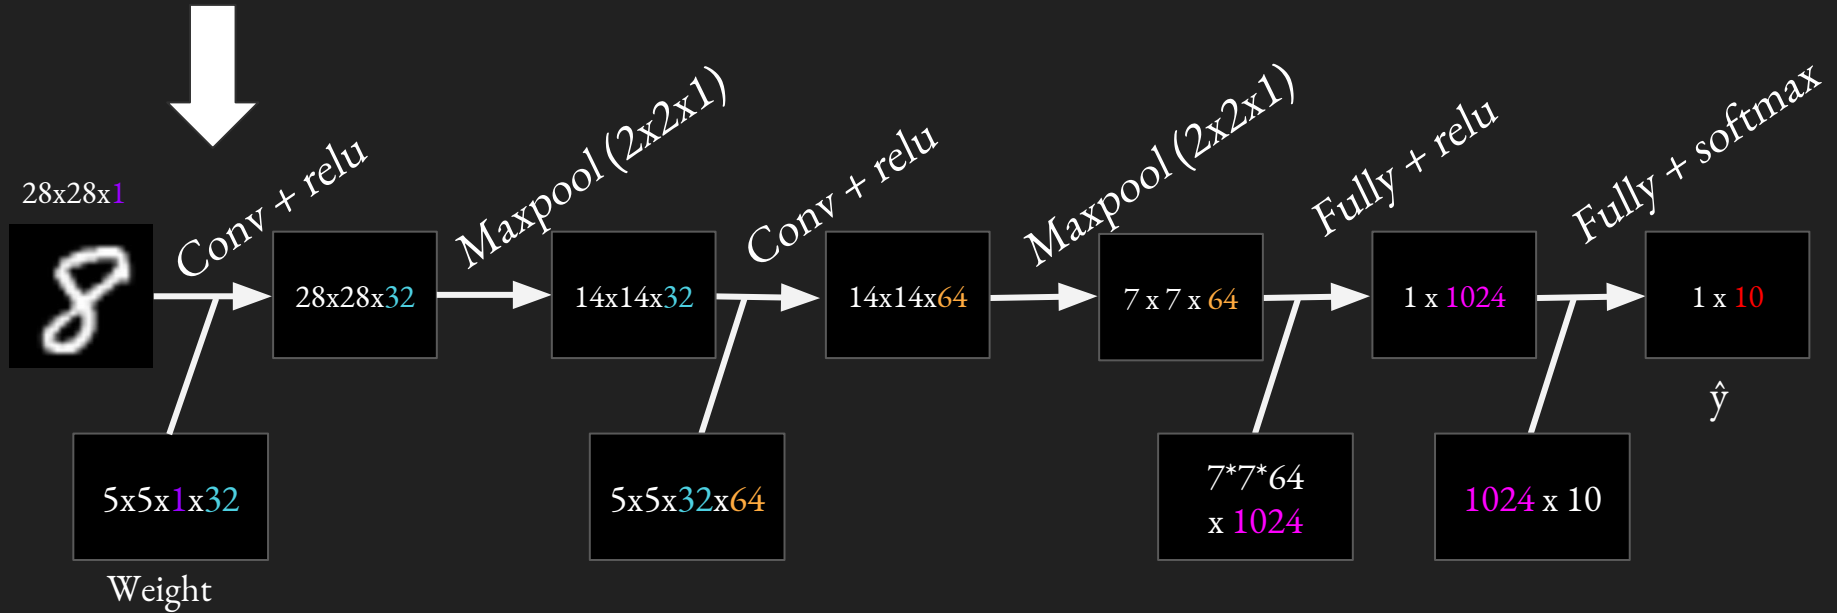
ConvNet with MNIST

Model



Strides for all convolutional layers: [1, 1, 1, 1]

Convolutional layer



```
conv = tf.nn.conv2d(images,  
                    kernel,  
                    strides=[1, 1, 1, 1],  
                    padding='SAME')
```

Convolutional layer: padding

"VALID" = without padding:

inputs:

	1	2	3	4	5	6	7	8	9	10	11	(12 13)
												dropped

"SAME" = with zero padding:

inputs:

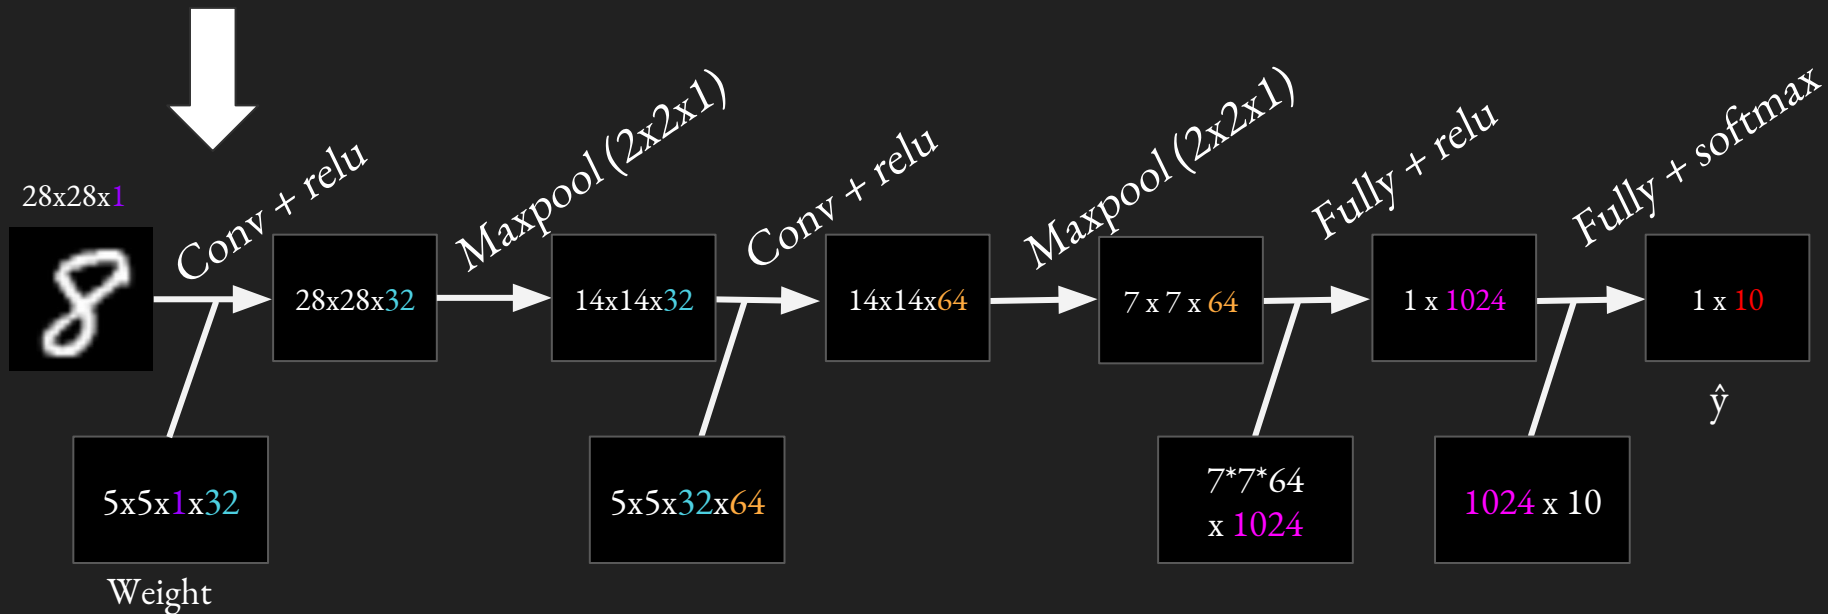
pad																		pad
	0	1	2	3	4	5	6	7	8	9	10	11	12	13	0	0		

Input width = 13

Filter width = 6

Stride = 5

Convolutional layer: Dimension

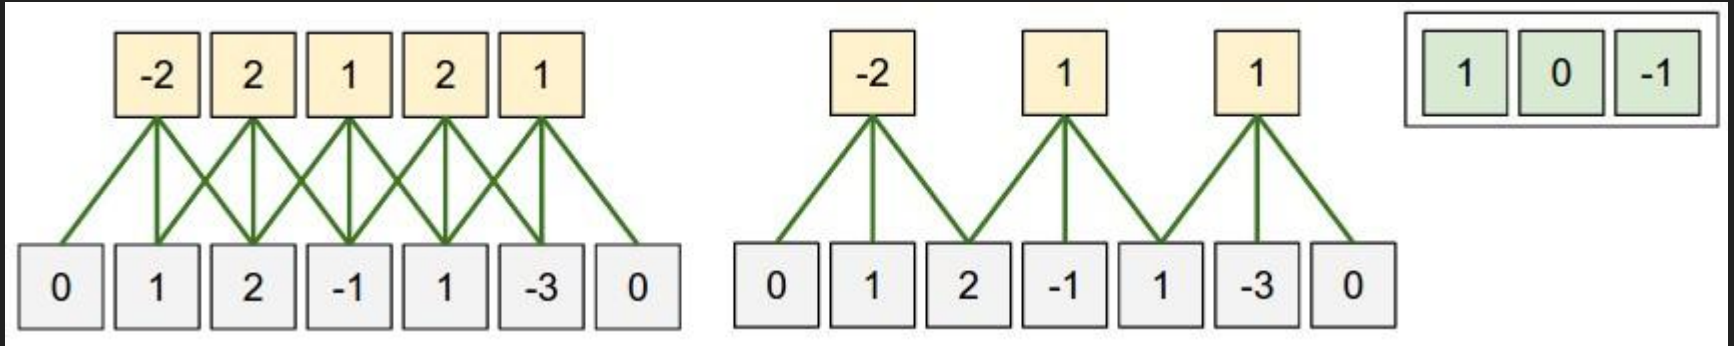


$$(\mathbf{W} - \mathbf{F} + 2\mathbf{P}) / \mathbf{S} + 1$$

W: input width/depth
P: padding

F: filter width/depth
S: stride

Convolutional layer: Dimension

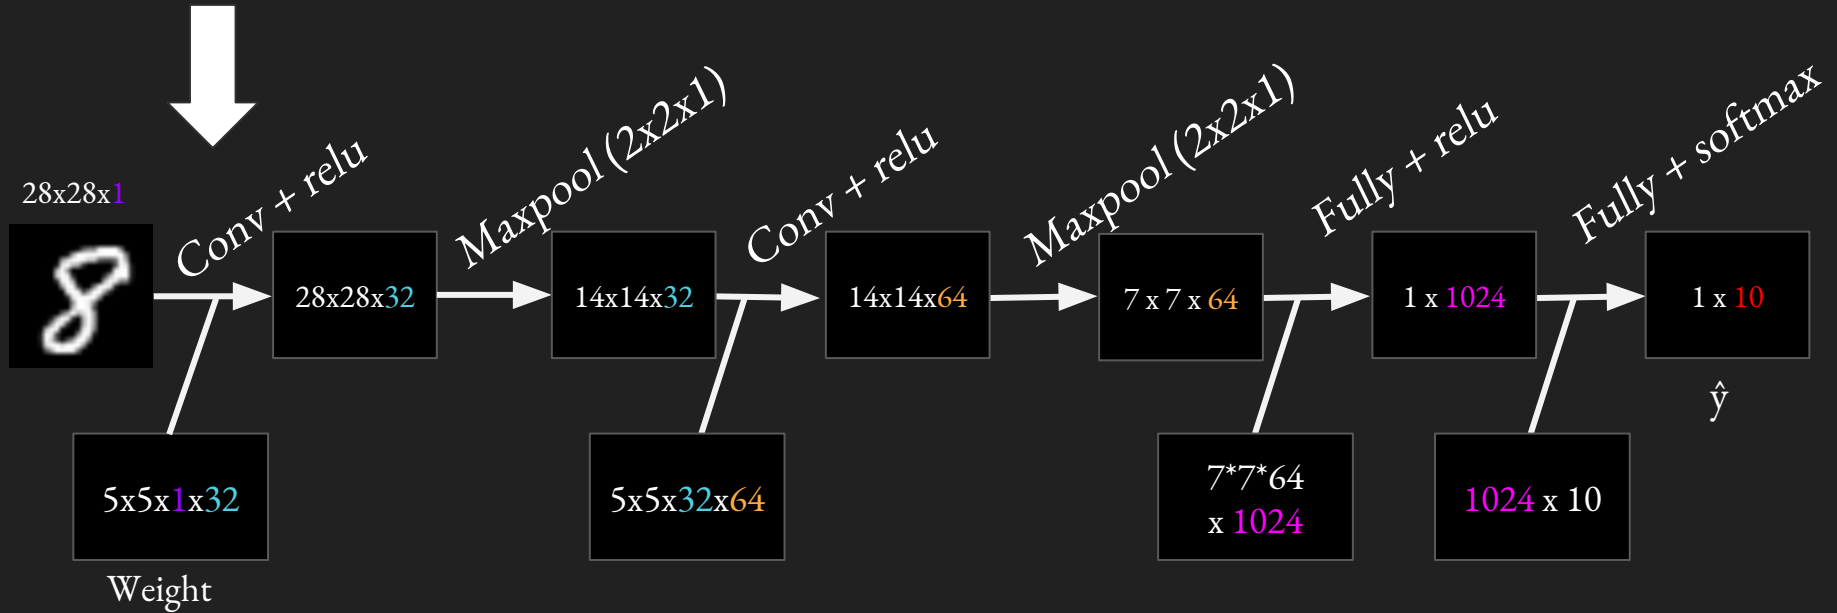


$$(W - F + 2P) / S + 1$$

W: input width/depth
P: padding

F: filter width/depth
S: stride

Convolutional layer: Dimension



$$(\mathbf{W}-\mathbf{F}+2\mathbf{P})/\mathbf{S}+1$$

$$(28-5+2*2)/1+1=28$$

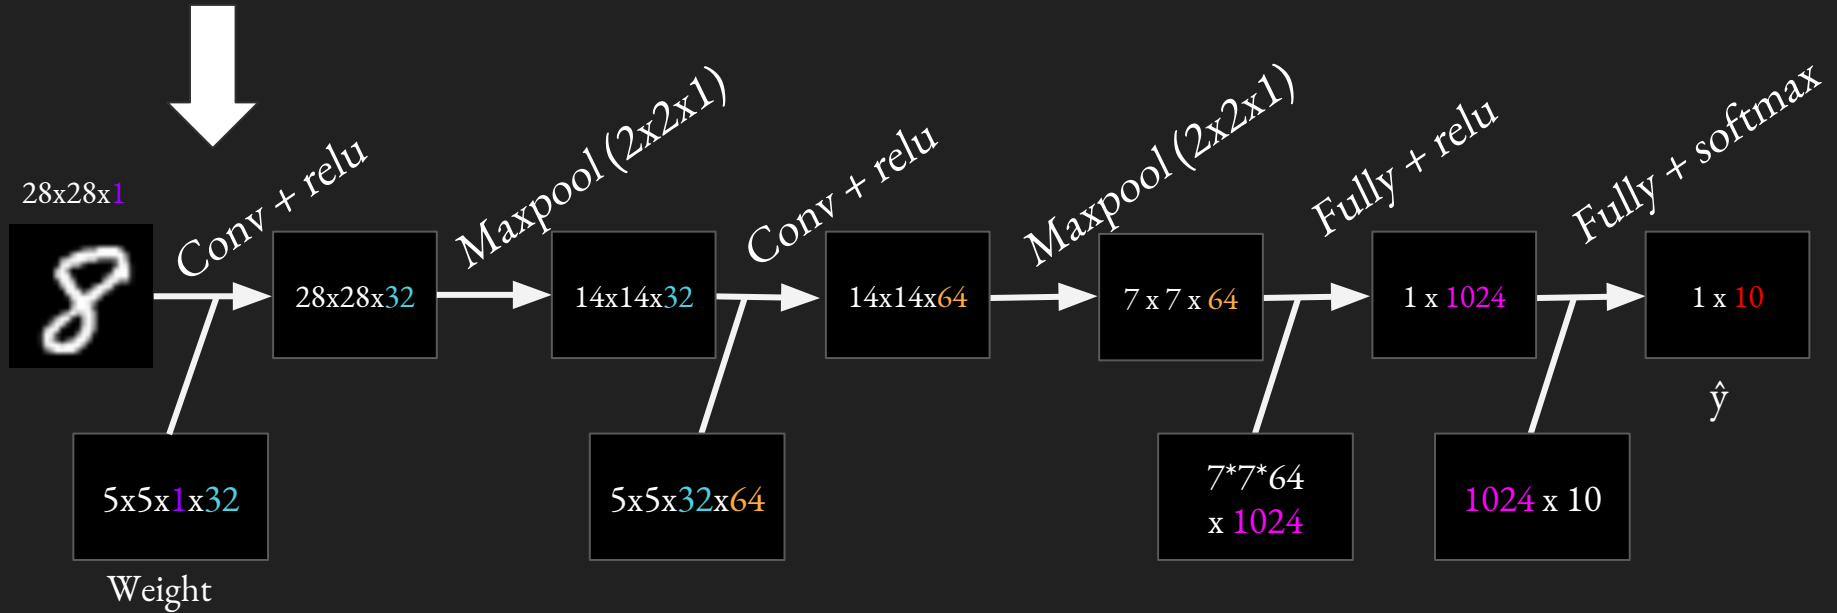
W: input width/depth

P: padding

F: filter width/depth

S: stride

Convolutional layer: Dimension



$$(\mathbf{W} - \mathbf{F} + 2\mathbf{P}) / \mathbf{S} + 1$$

$$(28 - 5 + 2 \cdot 2) / 1 + 1 = 28$$

W: input width/depth

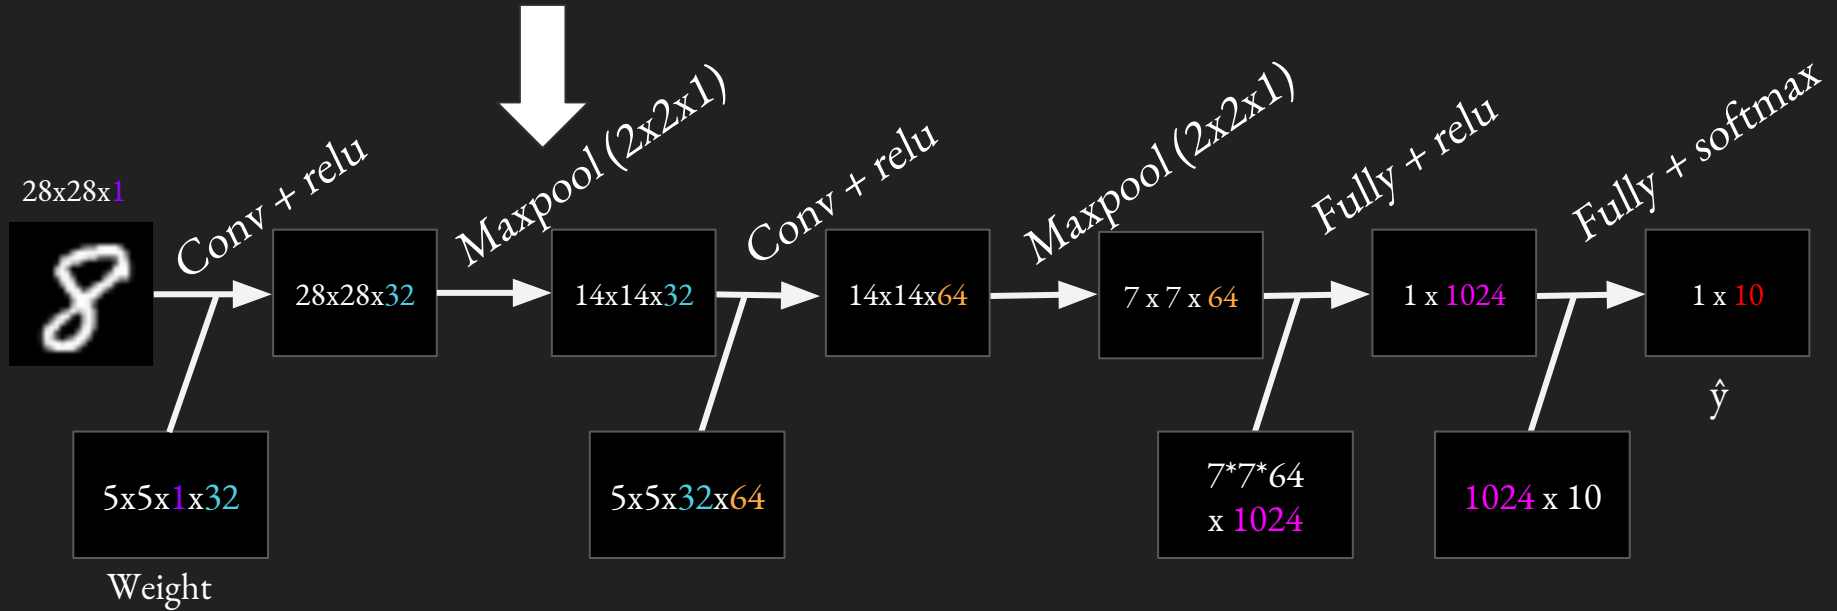
P: padding

F: filter width/depth

S: stride

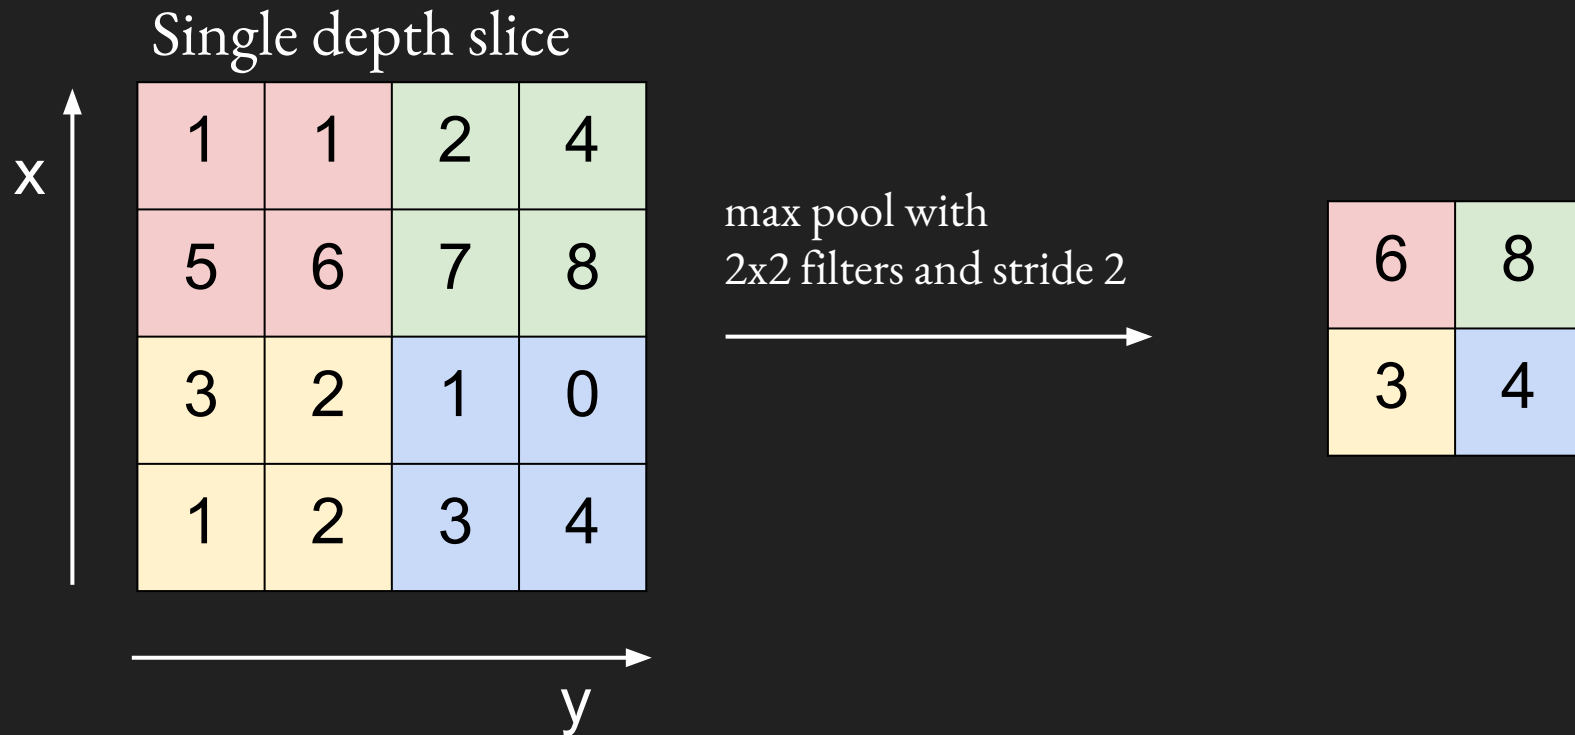
TF computes padding for us!

Maxpooling

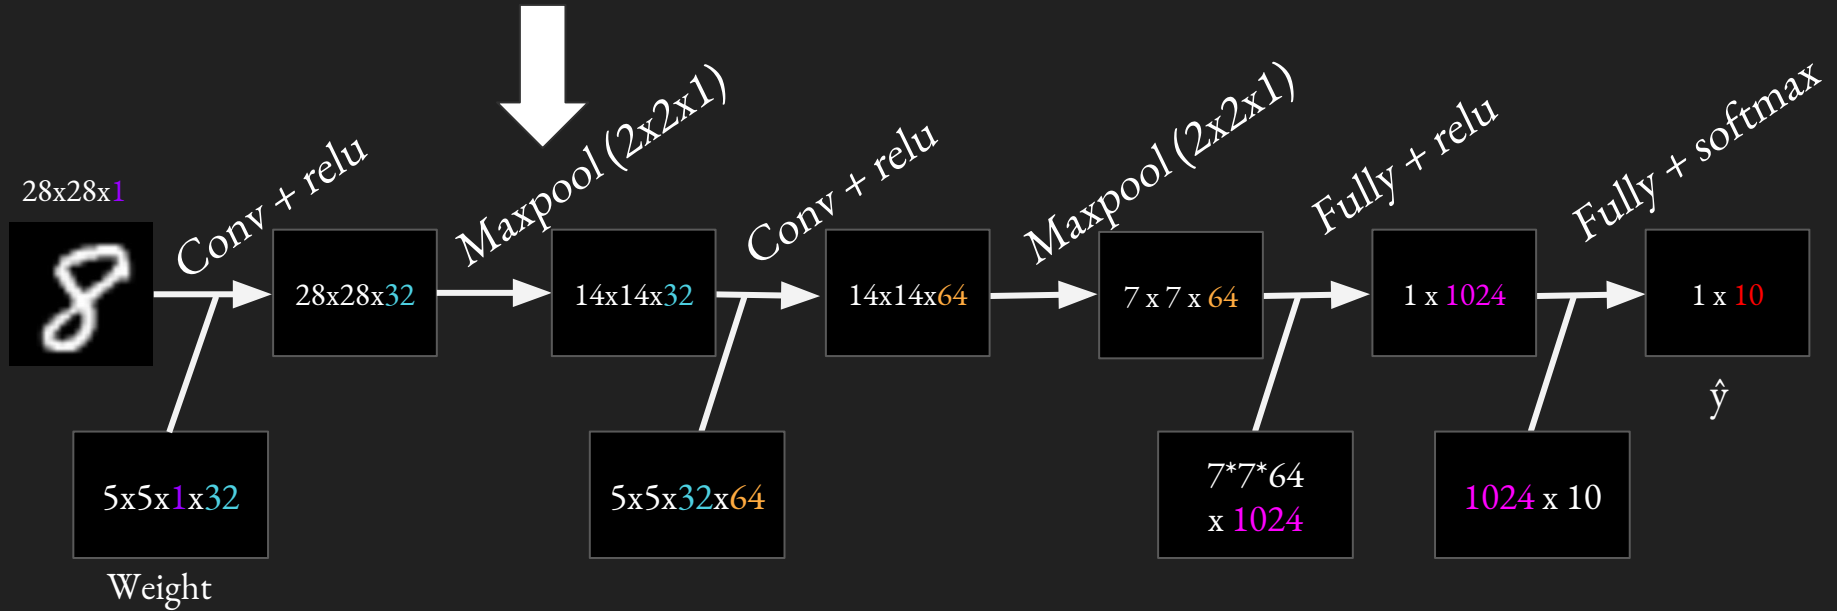


```
pool1 = tf.nn.max_pool(conv1,  
                        ksize=[1, 2, 2, 1],  
                        strides=[1, 2, 2, 1],  
                        padding='SAME')
```

Maxpooling



Maxpooling: Dimension

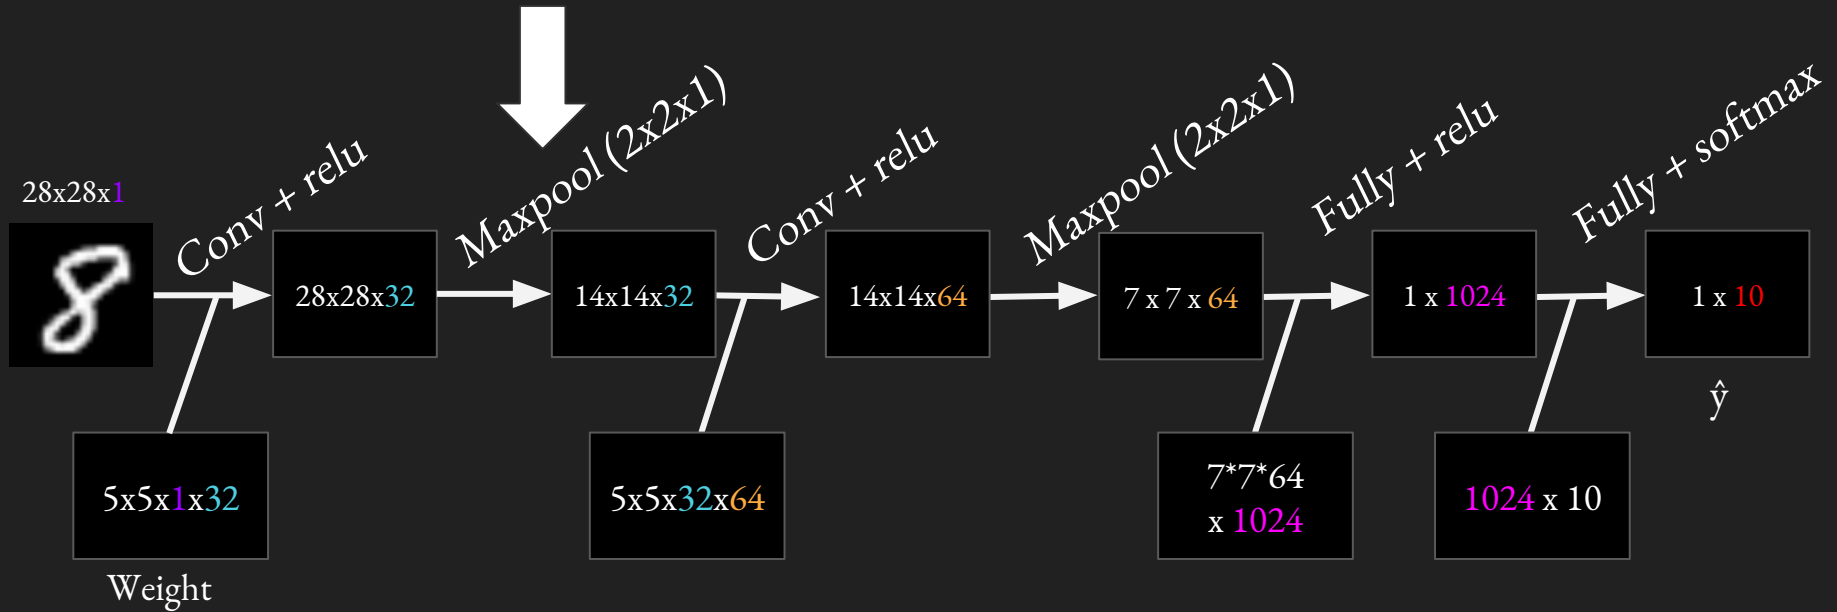


$$(\mathbf{W} - \mathbf{K} + 2\mathbf{P}) / \mathbf{S} + 1$$

W: input width/depth
P: padding

K: window width/depth
S: stride

Maxpooling: Dimension



$$(\mathbf{W} - \mathbf{K} + 2\mathbf{P}) / \mathbf{S} + 1$$

$$(28 - 2 + 2 \times 0) / 2 + 1 = 14$$

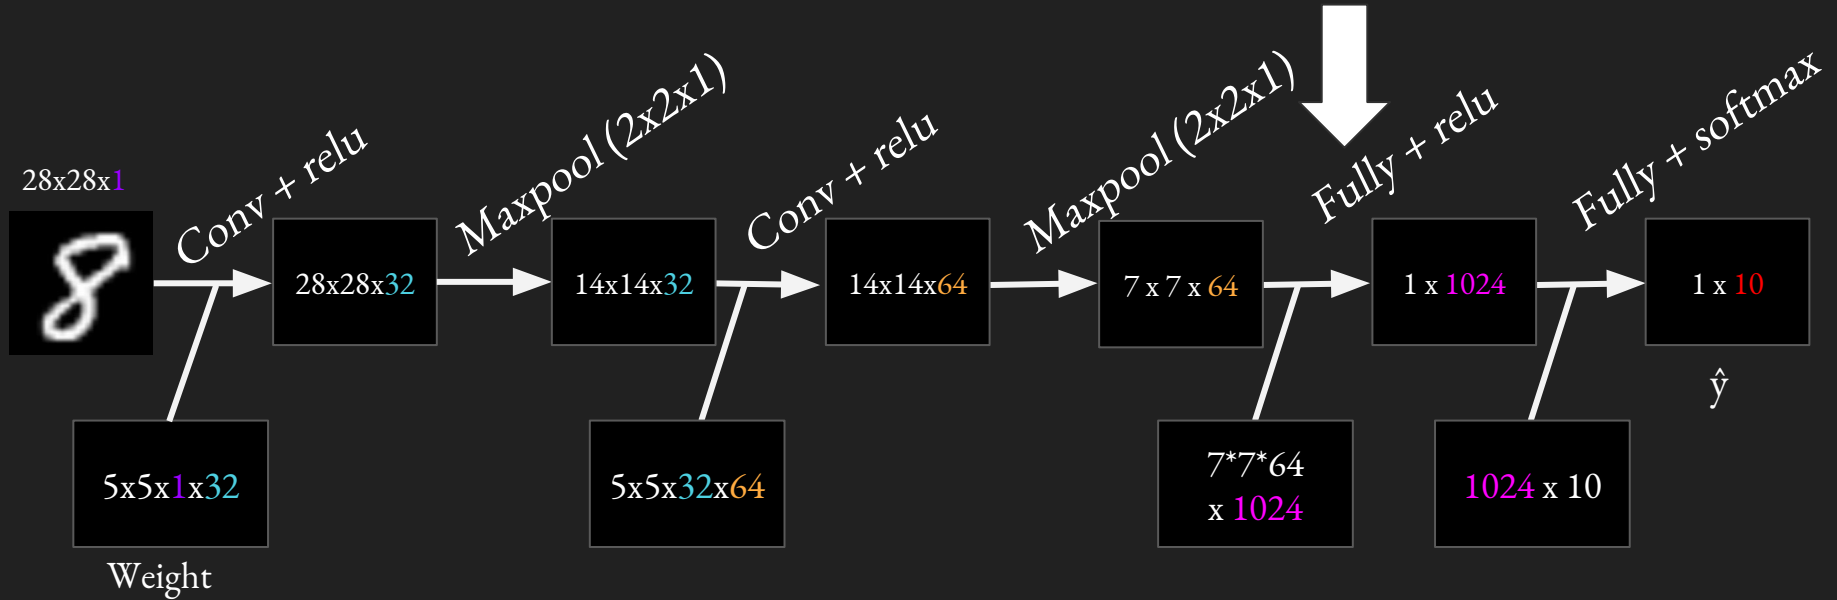
W: input width/depth

P: padding

K: window width/depth

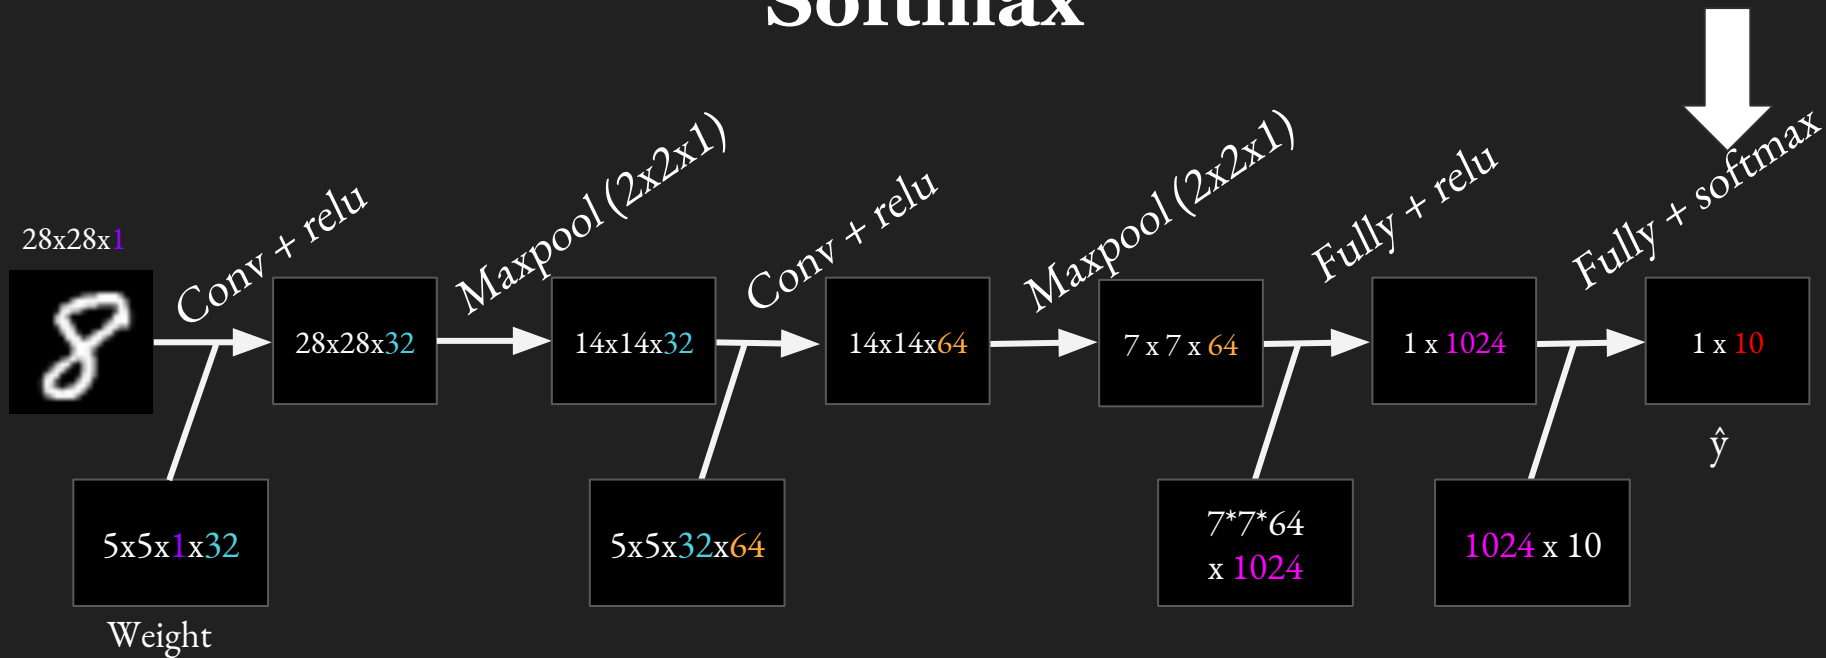
S: stride

Fully connected



$$fc = \text{tf.matmul}(\text{pool2}, w) + b$$

Softmax



Loss function

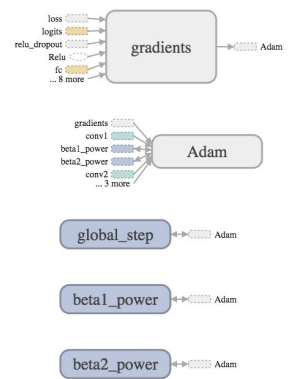
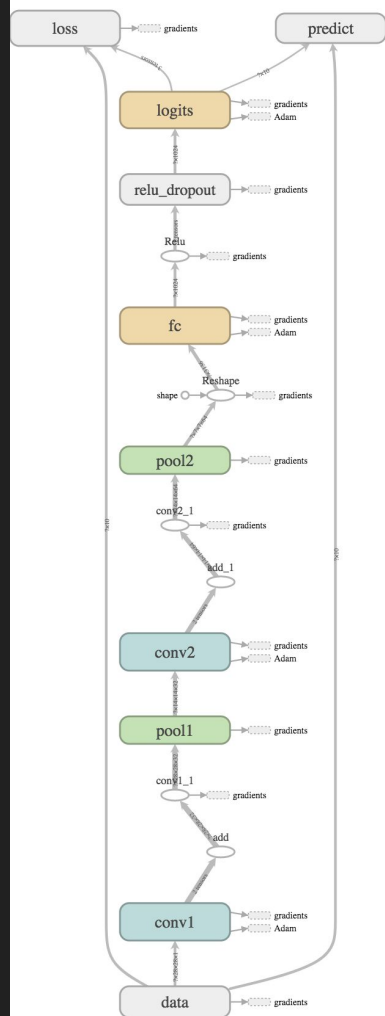
`tf.nn.softmax_cross_entropy_with_logits(labels=Y, logits=logits)`

Predict


`tf.nn.softmax(logits_batch)`

Interactive coding

07_convnet_mnist_starter.py from GitHub!
Update utils.py



Training progress



Test accuracy increases while training loss decreases!

Accuracy

Epochs	Accuracy
1	0.9131
2	0.9363
3	0.9478
5	0.9573
10	0.971
25	0.9818



tf.layers

tf.layers

We've been learning it the hard way

tf.layers.conv2d

```
conv1 = tf.layers.conv2d(inputs=self.img,  
                          filters=32,  
                          kernel_size=[5, 5],  
                          padding='SAME',  
                          activation=tf.nn.relu,  
                          name='conv1')
```

tf.layers.conv2d

```
conv1 = tf.layers.conv2d(inputs=self.img,  
                          filters=32,  
                          kernel_size=[5, 5],  
                          padding='SAME',  
                          activation=tf.nn.relu,  
                          name='conv1')
```

can choose
non-linearity to use

tf.layers.max_pooling2d

```
pool1 = tf.layers.max_pooling2d(inputs=conv1,  
                                pool_size=[2, 2],  
                                strides=2,  
                                name='pool1')
```


tf.layers.dense

```
fc = tf.layers.dense(pool2, 1024, activation=tf.nn.relu, name='fc')
```

tf.layers.dense

```
dropout = tf.layers.dropout(fc,  
                             self.keep_prob,  
                             training=self.training,  
                             name='dropout')
```

Drop neurals during training
Want to use all of them during testing

Next class

TFRecord

CIFAR

Style Transfer

Feedback: chiphuyen@cs.stanford.edu

Thanks!