

More LSH

LS Families of Hash Functions

LSH for Cosine Distance

Special Approaches for High Jaccard
Similarity

Jeffrey D. Ullman
Stanford University



Distance Measures

- Generalized LSH is based on some kind of “distance” between points.
 - Similar points are “close.”
- **Example:** Jaccard similarity is not a distance; 1 minus Jaccard similarity is.


Axioms of a Distance Measure

- d is a *distance measure* if it is a function from pairs of points to real numbers such that:
 1. $d(x,y) \geq 0$.
 2. $d(x,y) = 0$ iff $x = y$.
 3. $d(x,y) = d(y,x)$.
 4. $d(x,y) \leq d(x,z) + d(z,y)$ (*triangle inequality*).

Some Euclidean Distances

- L_2 norm: $d(x,y)$ = square root of the sum of the squares of the differences between x and y in each dimension.
 - The most common notion of “distance.”
- L_1 norm: sum of the differences in each dimension.
 - *Manhattan distance* = distance if you had to travel along coordinates only.

Examples of Euclidean Distances



Some Non-Euclidean Distances

- *Jaccard distance* for sets = 1 minus Jaccard similarity.
- *Cosine distance* for vectors = angle between the vectors.
- *Edit distance* for strings = number of inserts and deletes to change one string into another.

Example: Jaccard Distance

- Consider $x = \{1,2,3,4\}$ and $y = \{1,3,5\}$
- Size of intersection = 2; size of union = 5, Jaccard similarity (not distance) = $2/5$.
- $d(x,y) = 1 - (\text{Jaccard similarity}) = 3/5$.

Why J.D. Is a Distance Measure

- $d(x,y) \geq 0$ because $|x \cap y| \leq |x \cup y|$.
- $d(x,x) = 0$ because $x \cap x = x \cup x$.
 - And if $x \neq y$, then the size of $x \cap y$ is strictly less than the size of $x \cup y$.
- $d(x,y) = d(y,x)$ because union and intersection are symmetric.
- $d(x,y) \leq d(x,z) + d(z,y)$ trickier – next slide.

Triangle Inequality for J.D.

$$1 - \frac{|x \cap z|}{|x \cup z|} + 1 - \frac{|y \cap z|}{|y \cup z|} \geq 1 - \frac{|x \cap y|}{|x \cup y|}$$


Diagram illustrating the Triangle Inequality for Jaccard Distance (J.D.). The terms are circled in purple, and arrows indicate the corresponding distances:

- $d(x,z)$ points to $1 - \frac{|x \cap z|}{|x \cup z|}$
- $d(z,y)$ points to $1 - \frac{|y \cap z|}{|y \cup z|}$
- $d(x,y)$ points to $1 - \frac{|x \cap y|}{|x \cup y|}$

- **Remember:** $|a \cap b| / |a \cup b|$ = probability that $\text{minhash}(a) = \text{minhash}(b)$.
- Thus, $1 - |a \cap b| / |a \cup b|$ = probability that $\text{minhash}(a) \neq \text{minhash}(b)$.

Triangle Inequality – (2)

- **Claim:** $\text{prob}[\text{minhash}(x) \neq \text{minhash}(y)] \leq \text{prob}[\text{minhash}(x) \neq \text{minhash}(z)] + \text{prob}[\text{minhash}(z) \neq \text{minhash}(y)]$
- **Proof:** whenever $\text{minhash}(x) \neq \text{minhash}(y)$, at least one of $\text{minhash}(x) \neq \text{minhash}(z)$ and $\text{minhash}(z) \neq \text{minhash}(y)$ must be true.



$\text{minhash}(x) \neq \text{minhash}(z)$

$\text{minhash}(z) \neq \text{minhash}(y)$

Cosine Distance

- Think of a point as a vector from the origin $(0,0,\dots,0)$ to its location.
- Two points' vectors make an angle, whose cosine is the normalized dot-product of the vectors: $p_1 \cdot p_2 / |p_2| |p_1|$.
 - **Example:** $p_1 = 00111$; $p_2 = 10011$.
 - $p_1 \cdot p_2 = 2$; $|p_1| = |p_2| = \sqrt{3}$.
 - $\cos(\theta) = 2/3$; θ is about 48 degrees.

Edit Distance

- The *edit distance* of two strings is the number of inserts and deletes of characters needed to turn one into the other.
- An equivalent definition: $d(x,y) = |x| + |y| - 2|LCS(x,y)|$.
 - LCS = *longest common subsequence* = any longest string obtained both by deleting from x and deleting from y .

Example: Edit Distance

- $x = abcde$; $y = bcduve$.
- Turn x into y by deleting a , then inserting u and v after d .
 - Edit distance = 3.
- Or, computing edit distance through the LCS, note that $\text{LCS}(x,y) = bcde$.
- Then: $|x| + |y| - 2|\text{LCS}(x,y)| = 5 + 6 - 2*4 = 3$
= edit distance.

LSH Families of Hash Functions

Definition

Combining hash functions

Making steep S-Curves


Hash Functions Decide Equality

- There is a subtlety about what a “hash function” is, in the context of LSH families.
- A hash function h really takes two elements x and y , and returns a decision whether x and y are candidates for comparison.
- **Example**: the family of minhash functions computes minhash values and says “yes” iff they are the same.
- **Shorthand**: “ $h(x) = h(y)$ ” means h says “yes” for pair of elements x and y .

LSH Families Defined

- Suppose we have a space S of points with a distance measure d .
- A family \mathbf{H} of hash functions is said to be *(d_1, d_2, p_1, p_2) -sensitive* if for any x and y in S :
 1. If $d(x, y) \leq d_1$, then the probability over all h in \mathbf{H} , that $h(x) = h(y)$ is at least p_1 .
 2. If $d(x, y) \geq d_2$, then the probability over all h in \mathbf{H} , that $h(x) = h(y)$ is at most p_2 .

LS Families: Illustration



Example: LS Family

- Let:
 - S = subsets of some universal set,
 - d = Jaccard distance,
 - H formed from the minhash functions for all permutations of the universal set.
- Then $\text{Prob}[h(x)=h(y)] = 1-d(x,y)$.
 - Restates theorem about Jaccard similarity and minhashing in terms of Jaccard distance.

Example: LS Family – (2)

- **Claim:** H is a $(\boxed{1/3}, \boxed{3/4}, \boxed{2/3}, 1/4)$ -sensitive family for S and d .

If distance $\geq 3/4$
(so similarity $\leq 1/4$)

Then probability
that minhash values
agree is $\leq 1/4$

If distance $\leq 1/3$
(so similarity $\geq 2/3$)

Then probability
that minhash values
agree is $\geq 2/3$

For Jaccard similarity, minhashing gives us a $(d_1, d_2, (1-d_1), (1-d_2))$ -sensitive family for any $d_1 < d_2$.

Amplifying an LSH-Family

- The “bands” technique we learned for signature matrices carries over to this more general setting.
 - **Goal:** the “S-curve” effect seen there.
- AND construction like “rows in a band.”
- OR construction like “many bands.”

AND of Hash Functions

- Given family \mathbf{H} , construct family \mathbf{H}' whose members each consist of r functions from \mathbf{H} .
- For $h = \{h_1, \dots, h_r\}$ in \mathbf{H}' , $h(x)=h(y)$ if and only if $h_i(x)=h_i(y)$ for all i .
- **Theorem:** If \mathbf{H} is (d_1, d_2, p_1, p_2) -sensitive, then \mathbf{H}' is $(d_1, d_2, (p_1)^r, (p_2)^r)$ -sensitive.
 - **Proof:** Use fact that h_i 's are independent.

OR of Hash Functions

- Given family \mathbf{H} , construct family \mathbf{H}' whose members each consist of b functions from \mathbf{H} .
- For $h = \{h_1, \dots, h_b\}$ in \mathbf{H}' , $h(x)=h(y)$ if and only if $h_i(x)=h_i(y)$ for **some** i .
- **Theorem**: If \mathbf{H} is (d_1, d_2, p_1, p_2) -sensitive, then \mathbf{H}' is $(d_1, d_2, 1-(1-p_1)^b, 1-(1-p_2)^b)$ -sensitive.

Effect of AND and OR Constructions

- AND makes all probabilities shrink, but by choosing r correctly, we can make the lower probability approach 0 while the higher does not.
- OR makes all probabilities grow, but by choosing b correctly, we can make the upper probability approach 1 while the lower does not.

Composing Constructions

- As for the signature matrix, we can use the AND construction followed by the OR construction.
 - Or vice-versa.
 - Or any sequence of AND's and OR's alternating.

AND-OR Composition

- Each of the two probabilities p is transformed into $1-(1-p^r)^b$.
 - The “S-curve” studied before.
- **Example:** Take \mathbf{H} and construct \mathbf{H}' by the AND construction with $r = 4$. Then, from \mathbf{H}' , construct \mathbf{H}'' by the OR construction with $b = 4$.

Table for Function $1-(1-p^4)^4$

p	$1-(1-p^4)^4$
.2	.0064
.3	.0320
.4	.0985
.5	.2275
.6	.4260
.7	.6666
.8	.8785
.9	.9860

Example: Transforms a $(.2, .8, .8, .2)$ -sensitive family into a $(.2, .8, .8785, .0064)$ -sensitive family.

OR-AND Composition

- Each of the two probabilities p is transformed into $(1-(1-p)^b)^r$.
 - The same S-curve, mirrored horizontally and vertically.
- **Example:** Take \mathbf{H} and construct \mathbf{H}' by the OR construction with $b = 4$. Then, from \mathbf{H}' , construct \mathbf{H}'' by the AND construction with $r = 4$.

Table for Function $(1-(1-p)^4)^4$

p	$(1-(1-p)^4)^4$
.1	.0140
.2	.1215
.3	.3334
.4	.5740
.5	.7725
.6	.9015
.7	.9680
.8	.9936

Example: Transforms a $(.2, .8, .8, .2)$ -sensitive family into a $(.2, .8, .9936, .1215)$ -sensitive family.


Cascading Constructions

- **Example:** Apply the $(4,4)$ OR-AND construction followed by the $(4,4)$ AND-OR construction.
- Transforms a $(.2,.8,.8,.2)$ -sensitive family into a $(.2,.8,.9999996,.0008715)$ -sensitive family.

General Use of S-Curves

- For each AND-OR S-curve $1-(1-p^r)^b$, there is a *threshold* t , for which $1-(1-t^r)^b = t$.
- Above t , high probabilities are increased; below t , low probabilities are decreased.
- You improve the sensitivity as long as the low probability is less than t , and the high probability is greater than t .
 - Iterate as you like.
- Similar observation for the OR-AND type of S-curve: $(1-(1-p)^b)^r$.

Visualization of Threshold



An LSH Family for Cosine Distance

Random Hyperplanes
Sketches (Signatures)


Random Hyperplanes – (1)

- For cosine distance, there is a technique analogous to minhashing for generating a $(d_1, d_2, (1-d_1/180), (1-d_2/180))$ -sensitive family for any d_1 and d_2 .
- Called *random hyperplanes*.

Random Hyperplanes – (2)

- Each vector v determines a hash function h_v with two buckets.
- $h_v(x) = +1$ if $v \cdot x > 0$; $h_v(x) = -1$ if $v \cdot x < 0$.
- LS-family \mathbf{H} = set of all functions derived from any vector v .
- **Claim**: $\text{Prob}[h(x)=h(y)] = 1 - (\text{angle between } x \text{ and } y \text{ divided by } 180)$.

Proof of Claim



Signatures for Cosine Distance

- Pick some number of vectors, and hash your data for each vector.
- The result is a signature (*sketch*) of +1's and -1's that can be used for LSH like the minhash signatures for Jaccard distance.
- But you don't have to think this way.
- The existence of the LSH-family is sufficient for amplification by AND/OR.

Simplification

- We need not pick from among all possible vectors v to form a component of a sketch.
- It suffices to consider only vectors v consisting of $+1$ and -1 components.

Methods for High Degrees of Jaccard Similarity

Sets Represented by Sorted Strings
Use of String Length
Exploiting Prefixes

Setting: Sets as Strings

- We'll again talk about Jaccard similarity and distance of sets.
- However, now represent sets by strings (lists of symbols):
 1. Order the universal set.
 2. Represent a set by the string of its elements in sorted order.

Example: Shingles

- If the universal set is k-shingles, there is a natural lexicographic order.
- Think of each shingle as a single symbol.
- Then the 2-shingling of **abcad**, which is the set {ab, bc, ca, ad}, is represented by the list (*string*) [ab, ad, bc, ca] of length 4.

Example: Words

- If we treat a document as a set of words, we could order the words lexicographically.
- **Better:** Order words lowest-frequency-first.
- **Why?** We shall bucketize documents based on the early words in their lists.
 - Documents spread over more buckets.

Jaccard and Edit Distances

- Suppose two sets have Jaccard distance J and are represented by strings s_1 and s_2 . Let the LCS of s_1 and s_2 have length C and the (insert/delete) edit distance of s_1 and s_2 be E .
Then:

- $1-J = \text{Jaccard similarity} = C/(C+E)$.
- $J = E/(C+E)$.


Example: $s_1 = \text{acefh}$; $s_2 = \text{bcdegh}$.
LCS = ceh ; $C = 3$; $E = 5$; $1-J = 3/8$.

Works because these strings never repeat a symbol, and symbols appear in the same order.

Length-Based Indexes


- The simplest thing to do is create an index on the length of strings.
- A set whose string has length L can be Jaccard distance J from a set whose string has length M only if $L \times (1-J) \leq M \leq L / (1-J)$.
- **Example:** if $1-J = 90\%$ (Jaccard similarity), then M is between 90% and 111% of L .

Why the Limit on Lengths?



$$1-J \leq M/L$$
$$M \geq L \times (1-J)$$

A shortest candidate



$$1-J \leq L/M$$
$$M \leq L/(1-J)$$

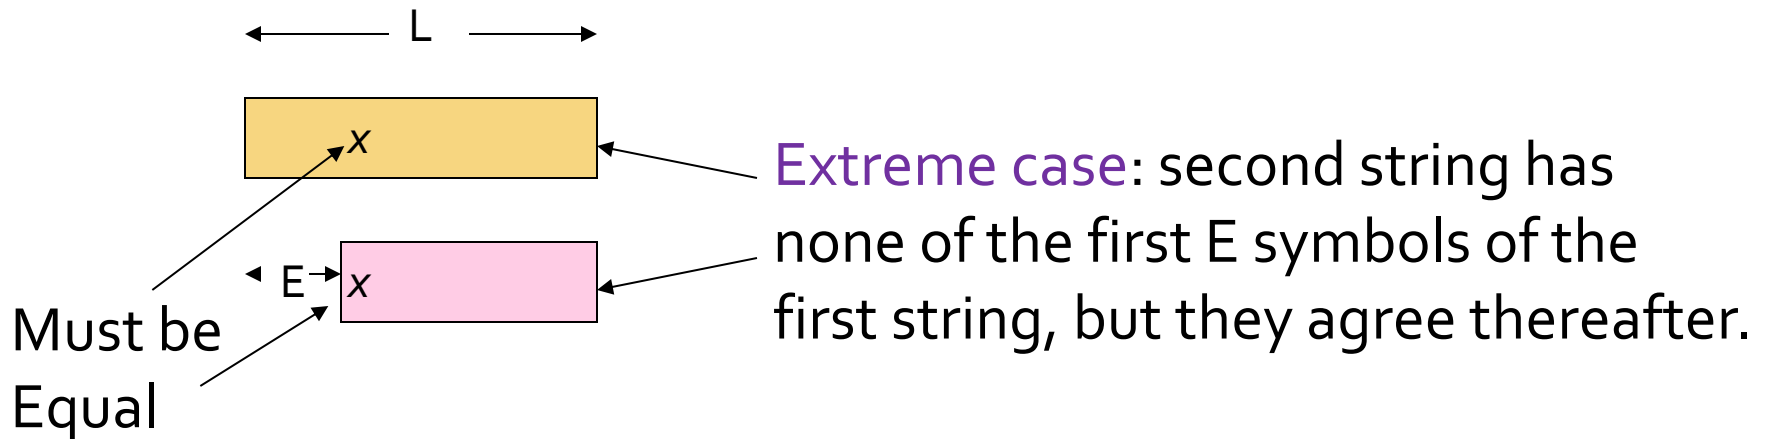
A longest candidate

Prefix-Based Indexing

- **Example:** If two strings are 90% similar, they must share some symbol in their *prefixes*.
 - These prefixes are of length just above 10% of the length of each string.
- **In general:** we can base an index on symbols in just the first $\lfloor JL+1 \rfloor$ positions of a string of length L .

Why the Limit on Prefixes?

Suppose a string of length L has E symbols
Before the first match with a second string.



If two strings do not share any of the first E symbols of the first string, then $J \geq E/L$.

Thus, $E = JL$ is possible, but any larger E is impossible. Index $E+1$ positions.

Indexing Prefixes

- Think of a bucket for each possible symbol.
- Each string of length L is placed in the bucket for the symbols in **each** of its first $\lfloor JL+1 \rfloor$ positions.

Lookup

- Given a *probe* string s of length L , with J the limit on Jaccard distance:

```
for (each symbol  $a$  among the  
    first  $\lfloor JL+1 \rfloor$  positions of  $s$ )  
    look for other strings in  
        the bucket for  $a$ ;
```


Example: Indexing Prefixes

- Let $J = 0.2$.
- String **abcdef** is indexed under a and b .
 - $\lfloor (0.2)^*6 + 1 \rfloor = 2$.
- String **acdfg** is indexed under a and c .
 - $\lfloor (0.2)^*5 + 1 \rfloor = 2$.
- String **bcde** is indexed only under b .
 - $\lfloor (0.2)^*4 + 1 \rfloor = 1$.
- If we search for strings similar to **cdef**, we need look only in the bucket for c .